

# Advanced Compile Features

*This presentation may contain forward-looking statements regarding Cadence's business or products. Actual results may differ materially from the information presented here.*

# Outline

**Parallel Netlist (PNL) for IXCOM flow**

**Parallel Invocation of HDL-ICE Analyzer**

**Parallel Synthesis**

**Parallel Partition Compiler (PPC)**


**Modular Compilation Flow (MC)**

# Parallel Netlist

- By default, HDL-ICE only generates a single netlist file for the entire design
- Parallel Netlist was introduced in VXE 19.10
- Compile with -parallelNL switch to the ixcom command:
  - `ixcom -parallelNL <other required ixcom options>`
- IXC.COM then generates multiple netlist files during HDL-ICE
  - Creates a netlist file for each module and a skeleton file for design hierarchy
- When Parallel netlist is detected, designImport is launched with multiple processes in parallel, and each process works on a subset of the netlist files

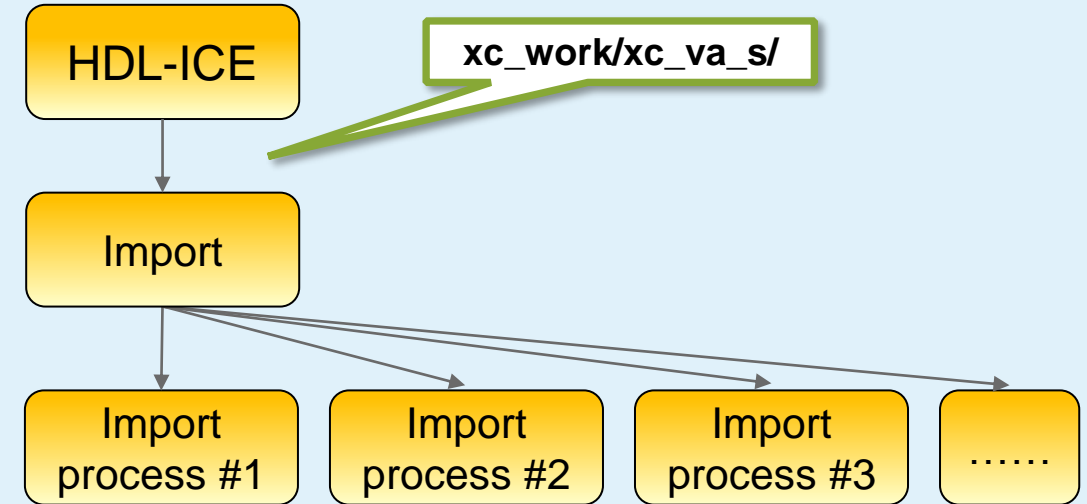
# Without Parallel NL v.s. with Parallel NL (PNL)

## Without PNL



- Single netlist for the entire design
- Single netlist file is processed sequentially
- Performance bottleneck when netlist grows bigger

## With PNL



- xc\_va.ver is replaced with a xc\_va\_s directory which contains:
  - Netlist files for all design modules: 1.v, 2.v, 3.v, etc.
  - A design skeleton file
- Multiple designImport processes are launched and each process works on a subset of the netlist files

# Without Parallel NL v.s. with Parallel NL (PNL)

## Without PNL

xc\_work/qel/ixcomImport.qel contents

```
refLib qtref generic
# design -rm
importOption {mode full}
importOption {special DELETE_LIBS_ON_FULL}
importOption {library hdlldb}
importOption {V1995 ON}
importOption {IXCOM RTL ON}
netlistFile {verilog xc_work/xc_va.ver}
importOption {toplib hdlldb} {topcell xcva_top}
designImport
```

Vaelab command in ixcom.log

```
vaelab -append_log -logfile xc_work/va.log -valib xc_work/va.lib -F xc_work/va_reflib_elab -DontStopOnError
-Work xc_ncwork -UCDBnoZZ -keepRtlSymbol -keepAllFlipFlop -outputVlog xc_work/xc_va.ver xcva_top
```

## With PNL

xc\_work/qel/ixcomImport.qel contents

```
refLib qtref generic
# design -rm
importOption {mode full}
importOption {special DELETE_LIBS_ON_FULL}
importOption {library hdlldb}
importOption {V1995 ON}
importOption {IXCOM RTL ON}
netlistFile {skeleton xc_work/xc_va.s}
importOption {toplib hdlldb} {topcell xcva_top}
designImport
```

Vaelab command in ixcom.log

```
vaelab -append_log -logfile xc_work/va.log -valib xc_work/va.lib -F xc_work/va_reflib_elab -DontStopOnError
-Work xc_ncwork -UCDBnoZZ -keepRtlSymbol -keepAllFlipFlop -outputSkeleton xc_work/xc_va.s xcva_top
```



# Without Parallel NL v.s. with Parallel NL (PNL)

## Without PNL

xe.msg snippet

```
import Elapsed time: 0:42:47, CPU time: 0:41:53.22, Memory Usage: 82.0G
```

- DesignImport launches a single process

## With PNL

xe.msg snippet

```
import Elapsed time: 0:00:02, CPU time: 0:00:00.38, Memory Usage: 3.1G
import Elapsed time: 0:00:05, CPU time: 0:00:00.41, Memory Usage: 3.1G
import Elapsed time: 0:01:07, CPU time: 0:00:57.86, Memory Usage: 5.2G
import Elapsed time: 0:03:09, CPU time: 0:02:53.10, Memory Usage: 8.9G
import Elapsed time: 0:03:23, CPU time: 0:02:47.22, Memory Usage: 8.9G
import Elapsed time: 0:03:30, CPU time: 0:02:57.55, Memory Usage: 9.0G
import Elapsed time: 0:03:34, CPU time: 0:03:03.06, Memory Usage: 9.0G
import Elapsed time: 0:03:35, CPU time: 0:03:02.36, Memory Usage: 9.0G
import Elapsed time: 0:03:37, CPU time: 0:02:57.38, Memory Usage: 9.0G
import Elapsed time: 0:03:41, CPU time: 0:03:03.25, Memory Usage: 8.9G
import Elapsed time: 0:03:44, CPU time: 0:02:57.91, Memory Usage: 8.9G
import Elapsed time: 0:03:51, CPU time: 0:03:06.36, Memory Usage: 8.9G
import Elapsed time: 0:10:50, CPU time: 0:07:27.16, Memory Usage: 21.6G
import Elapsed time: 0:10:59, CPU time: 0:07:36.92, Memory Usage: 20.1G
import Elapsed time: 0:11:43, CPU time: 0:02:32.91, Memory Usage: 10.6G
import Elapsed time: 0:11:55, CPU time: 0:00:01.15, Memory Usage: 4.2G
```

- DesignImport launches multiple processes in parallel
- Total elapsed time is equal to the largest elapsed time

# Parallel Netlist: Module Skeleton File

- The module skeleton file consists of:
  - Module NL definition without design contents (ports, nets, Q\_cells)
  - Module instantiation with module name and instance name only

```
// 1.v
module ASSERTION;
//pragma CVASTRPROP MODULE SKELETON "0.000000"
endmodule

// 2.v
module Q_RBUFZP;
//pragma CVASTRPROP MODULE SKELETON "1.000000"
endmodule

// 3.v
module Q_RBUFZN;
//pragma CVASTRPROP MODULE SKELETON "1.000000"
endmodule

// 4.v
module xc_top_1;
//pragma CVASTRPROP MODULE SKELETON "8263.000000"
Q_RBUFZP dum2 ();
Q_RBUFZN dum1 ();
endmodule
```

# Parallel Netlist: Performance

- While using Parallel NL feature, designImport elapsed time is improved by 2x to 3x

#	IXCOM based design	With/Without parallel NL (PNL)	Preoptimized gate count (million gates)	Overall IXCOM compile elapsed time (minutes)	designImport elapsed time (minutes)
1	18-board design	Without PNL	674	64	19
		With PNL		57	9
2	36-board design	Without PNL	1422	142	40
		With PNL		116	12



# Parallel Netlist: Guidelines & Limitations in VXE19.10

- Parallel NL doesn't support below use models:
  - Incremental compilation. User has to do a clean recompile every time
  - Encrypted designs
  - LEC flow
  - External netlist flow (ixcom -externalNetlists)
- VXE19.10 Parallel NL feature supports IXCOM RTL flow only

# Outline

**Parallel Netlist (PNL) for IXCOM flow**

**Parallel Invocation of HDL-ICE Analyzer**

**Parallel Synthesis**

**Parallel Partition Compiler (PPC)**

**Modular Compilation Flow (MC)**

# Parallel Invocation of HDL-ICE Analyzer: Introduction

- The HDL-ICE analysis sessions might take significant time for designs that have multiple libraries. Run these sessions in parallel can optimize the compilation time.
- IXCOM generates the vavlog and vavhdl commands for each library and merges all the sessions of one library in a single session.
  - Use IXCOM option `-parallelHdliceAnalyze [<nJobs>]`
- When `-parallelHdliceAnalyze [<nJobs>]` is used, IXCOM automatically invokes the vavlog commands in parallel using the following command:
  - `make -j <nJobs> -f xc_work/xc_make vavlog_parallel`
  - `<nJobs>` is the number of vavlog sessions required by user and is optional.
    - If it is not specified, IXCOM compares the number of the available cores on the machine and the number of vavlog sessions that are required, and sets the lesser number as `<nJobs>`.
    - If it is specified, but its value exceeds the number of available cores on the machine, the number of available cores on machine is set as `<nJobs>`.

# Parallel Invocation of HDL-ICE Analyzer: Example

- Upon executing `-parallelHdliceAnalyze`, IXCOM generates a target rule, `vavlog_parallel`, in `xc_work/xc_make` Makefile. For example:

```
vavlog_parallel: vavlog_1 vavlog_2 vavlog_3 vavlog_4 vavlog_5
```

- It also creates a target rule for each library in `xc_work/xc_make` as follows:

```
vavlog_1:  
vavlog $(VAVLOG_OPTIONS) -append_log $(VALIB) +ixcom -Work IXCOM_TEMP_LIBRARY \  
-File xc_work/v/vavlog-IXCOM_TEMP_LIBRARY-verilog.f  
vavlog_2:  
vavlog $(VAVLOG_OPTIONS) -append_log $(VALIB) +ixcom -Work lib1 \  
-File xc_work/v/vavlog-lib1-verilog.f  
vavlog_3:  
vavlog $(VAVLOG_OPTIONS) -append_log $(VALIB) +ixcom -Work lib2 \  
-File xc_work/v/vavlog-lib2-verilog.f  
vavlog_4:  
vavlog $(VAVLOG_OPTIONS) -append_log $(VALIB) +ixcom -Work lib3 \  
-File xc_work/v/vavlog-lib3-verilog.f  
vavlog_5:  
vavlog $(VAVLOG_OPTIONS) -append_log $(VALIB) +ixcom -Work lib4 \  
-File xc_work/v/vavlog-lib4-verilog.f
```

# Parallel Invocation of HDL-ICE Analyzer: Example (Con'd)

- Then, IXCOM invokes HDL-ICE Compiler as follows:

```
make -f xc_work/xc_make vlog_serial  
make -f xc_work/xc_make vlog_parallel -j 5  
make -f xc_work/xc_make vaelab_only
```

- In the ixcom.log file, an INFO message is generated to indicate the number of job slots used for the parallel invocation:

```
Info!      Invoking hdllice analysis steps in parallel using (5) job slots.
```

# Outline

**Parallel Netlist (PNL) for IXCOM flow**

**Parallel Invocation of HDL-ICE Analyzer**

**Parallel Synthesis**

**Parallel Partition Compiler (PPC)**

**Modular Compilation Flow (MC)**



# Parallel Synthesis: Introduction

- Parallel synthesis is useful for synthesizing VHDL or Verilog RTL designs containing tens of million gates.
  - Enables distribution of synthesis into tasks at module level. The tasks are dynamically allocated to balance the load across multiple workstations or processes.
  - Helps to increase the speed during design synthesis.
- Ways to enable parallel synthesis
  - Using parallel.host file
  - Using multiple CPUs of the same workstation
  - Using Load Sharing Facility (LSF) queue
  - Using Network Computer (NC)

# Parallel Synthesis: Using parallel.host File

- Hosts and job numbers need to be specified in parallel.host file:

```
HostA -jobs=20  
HostB -jobs=30
```

- 20 HDL-ICE processes are specified to start on HostA and 30 to start on HostB
- The maximum number of total parallel HDL-ICE processes on all hosts is 248
- File parallel.host should be located in design directory
- The workstations can be of mixed platforms
- Execute any of the following commands to start parallelsynthesis using the parallel.host file:
  - vaelab -parallel <top\_module>
  - hdlSynthesize -parallel <top\_module>
  - vaelab -parallelHostSSH <top\_module>
  - hdlSynthesize -parallelHostSSH <top\_module>
    - With -parallelHostSSH, HDL-ICE uses SSH to distribute synthesis over hosts listed in parallel.host.

# Parallel Synthesis: Using Multiple CPUs of the Same Workstation

- If a workstation has multiple CPUs, tasks can be distributed to run in parallel on those CPUs.
- To distribute a definite number of tasks on the current workstation, use either of the following two commands:
  - `vaelab -parallel <number> <top_module>`
  - `hdlSynthesize -parallel<number> <top_module>`
    - The maximum <number> of parallel HDL-ICE Compiler processes on the current workstation is 248.

# Parallel Synthesis: Using Load Sharing Facility (LSF) Queue

- Synthesis processes can also be submitted to an LSF queue using any of the following commands:
  - `vaelab -parallelLSFSameHost <number> <top_module>`
  - `hdlSynthesize -parallelLSFSameHost<number> <top_module>`
  - `vaelab -parallelLSF <number> <top_module>`
  - `hdlSynthesize -parallelLSF<number> <top_module>`
    - `<number>` is a required argument that specify the number of jobs to submit to the LSF queue. The maximum `<number>` is 248.
    - Use `-parallelLSFSameHost` option to assign one workstation for HDL-ICE Compiler and avoid network latency between multi-workstations during parallel synthesis.

# Parallel Synthesis: Using Load Sharing Facility (LSF) Queue

## Using Multiple bsub Commands

- Starting HDLICE20.12, it supports the distribution of synthesis process over multiple hosts in the LSF network using multiple bsub commands.
  - One bsub command starts one slave process.
- The usage is as follows:
  - `vaelab -parallelLSFMultiBsub <number>`
  - `hdlSynthesize -parallelLSFMultiBsub<number>`
    - The maximum <number> is 500.
- Options to the bsub executable can be specified in a file called LSF\_options. The compiler does not check for syntax of these options.

# Parallel Synthesis: Using Network Computer (NC)

- If NC is used instead of LSF, parallel synthesis processes can also be submitted to NC using either of the following two commands:
  - `vaelab -parallelNCSameHost <number> <top_module>`
  - `hdlSynthesize -parallelNCSameHost<number> <top_module>`
    - Use `-parallelNCSameHost` option to assign one workstation to run parallel synthesis.
    - The maximum `<number>` is 248.



# Parallel Synthesis: Using Network Computer (NC)

## Using Multiple nc run Commands

- Starting HDLICE20.12, it supports the distribution of synthesis process over multiple hosts in the NC network using multiple nc run commands.
  - One nc run starts one slave process.
- The usage is as follows:
  - `vaelab -parallelMultiNC <number>`
  - `hdlSynthesize -parallelMultiNC<number>`
    - The maximum <number> is 500.
- If a file named NC\_options exist under the current working directory, the string will be appended to the NC command when a job is submitted. All the NC options must be defined in one line. The compiler does not check for syntax of these options.

# Outline

**Parallel Netlist (PNL) for IXCOM flow**

**Parallel Invocation of HDL-ICE Analyzer**

**Parallel Synthesis**

**Parallel Partition Compiler (PPC)**

**Modular Compilation Flow (MC)**

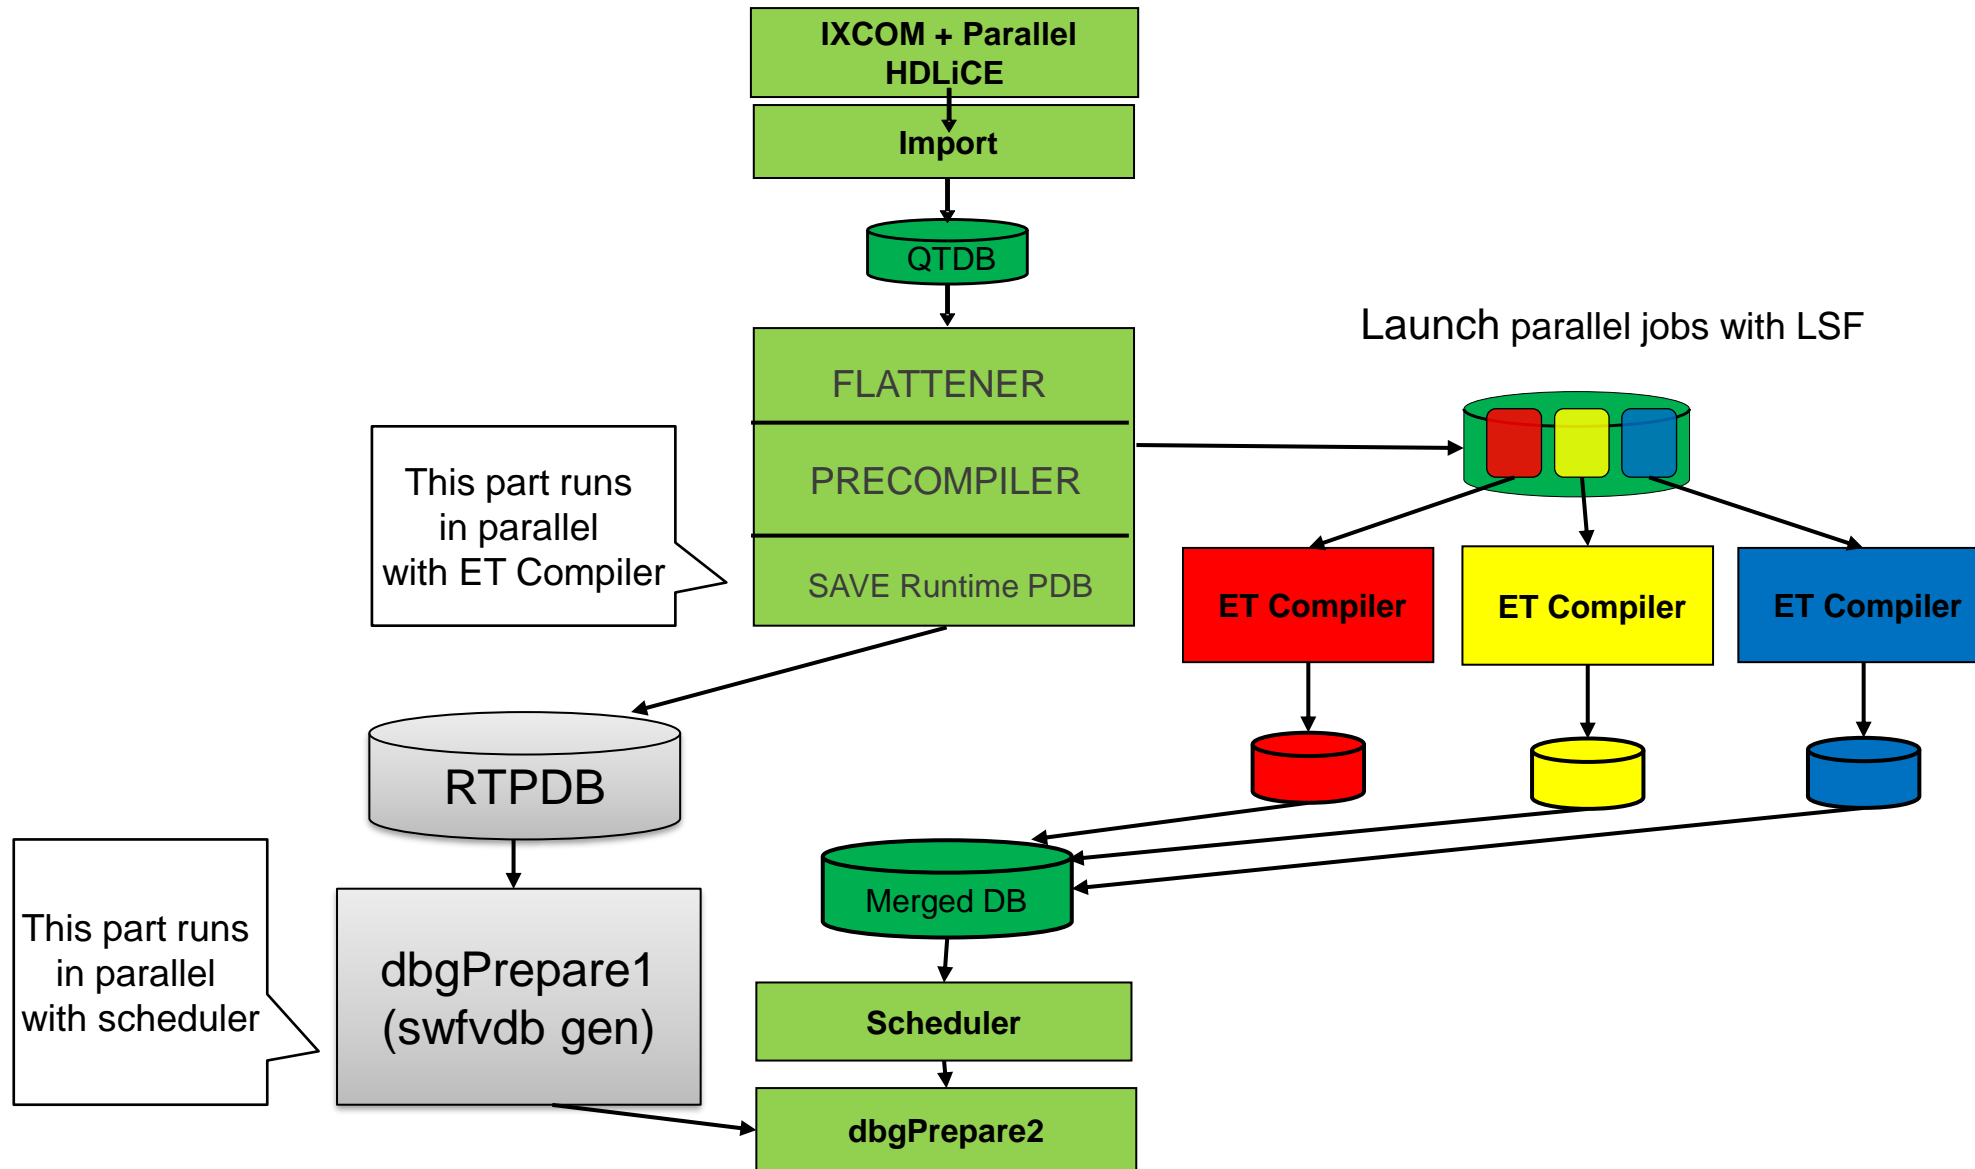
# Parallel Partition Compiler (PPC)

- By default, et3compile compiles the entire design sequentially
- VXE 181/185 introduced a new feature called parallel partition compiler (PPC)
- Motivations for PPC technology
  - Reduce compile time by focusing on the longest compile stage: et3compile. The et3compile time for larger designs is linear without the parallel process
- With PPC, the design is split into separate partitions at the end of precompile
- Most of et3compile's work is done on the partitions in parallel. Results are then merged and the design is scheduled.

# PPC Requirements

- Requires installation of workload management platform
  - Currently support LSF and NC
  - Support for others done with wrappers
    - PPC-invoked LSF commands then execute wrappers which call job scheduler commands
- Workstation memory requirements based on design size
  - 1GB memory per 1 million gates
  - Require multiple smaller memory machines for partition compile with LSF
  - PPC does not reduce host memory footprint compared to without PPC
- Sufficient disk space

# PPC Compilation Flow



# PPC: What It Is Doing

- Parallel Partition Compiler parallelizes the most time consuming section of the compile process
  - Design needs to be sliced into multi-board sized partitions
  - The partitions should be roughly the same size to improve the compile time
    - Variation in partitions will result in different compile time, affect overall compile times
  - Partitions are assigned boards and domains by user or by compiler
  - After partitions defined and domains assigned, normal compile with compile or compileFind command
  - Compile will spawn jobs for partitions automatically using LSF / NC
  - When spawned jobs finish, performs final steps including scheduling



# How to Enable PPC: Partitioning the Design (Step 1)

- Step 1: partition the design by one of the three ways
  - Manual partitioning (VXE186 or later)
    - Original use model
  - Automatic partitioning (VXE186 or later)
    - XEL user data command **autoPart**
  - Advanced automatic partitioning (VXE19.10 or later)
    - Second generation automatic partitioning
    - Original autoPart is still supported but no new development in VXE19.10
    - **compilerOption -add {advancedPart ON}**

# How to Enable PPC: Manual Partitioning

- First, use partitionGroup command to assign an instance to a selected partition:
  - `partitionGroup -add {<instance>.[]<prer>[*]<partition>}`
  - <instance> is either a wildcard character (\*) or a valid hierarchical instance name. <partition> is one of the partition identifiers part1- part65499.
  - Wildcard used in <instance> means everything not explicitly assigned
  - All instances not covered by specification are placed in lowest number partition

```
partitionGroup -add {dut.a.b part1}  
partitionGroup -add {dut.a.b.r* part2}  
partitionGroup -add {dut.a.c.* part3}  
partitionGroup -add {* part4}
```

- Secondly, use partitionAssign command to control the emulator configuration (boards and domains) assigned to a given partition.
  - `partitionAssign -add {<partition> <resource>}`

```
partitionAssign -add {part1 0-1}  
partitionAssign -add {part2 2-3}  
partitionAssign -add {part3 4}  
partitionAssign -add {part4 5}
```

# How to Enable PPC: autoPart

- XEL command autoPart automatically partitions and assigns resources for PPC
  - Partitioning is done early in precompile stage
  - autoPart overwrites any existing partitionGroup or partitionAssign command
  - It identifies suitable number of partitions of the design based on the resources and design status, and tries to group them into 3 or 4 board partitions
  - Solution generated is written into partitionGroup.qel, partitionAssign.qel, and tmp directory
- The generated files can be modified and used in future compile
  - autoPart generates an initial solution that user can modify with
  - User with design knowledge is key, understanding interactions of design
    - Move hierarchy in groupings around to improve IO
    - Modify the assignments by add / subtract domains to help utilization
    - Modify the assignments by altering the location of the domains on the rack

# How to Enable PPC: Advanced Automatic Partitioning

- In VXE 19.10 or later, advanced automatic partitioning can be turned on by
  - `compilerOption -add {advancedPart ON | <N>}`
  - ON option enables the compiler to create an appropriate number of partitions based on the boards/domains available
  - N is the number of partitions that can be specified to overwrite compiler's decision
  - Recommend Use Model:
    - *Complete compile first, then set partition number if partitions have high utilization*
    - Fewer, larger partitions are better to help reduce partitions utilization

# How to Enable PPC: Advanced Automatic Partitioning

- advancedPart is the second generation automatic partitioning
  - Original autoPart is still supported but no new development
  - Unlike autoPart, partitioning is done after precompile
  - Netlist contains all precompiler instrumentation logic
    - Works better with 1x and other features where the compiler instrumentation is done
  - Generates better fCLK and better design utilization
  - Partitions are generated by gate count
    - No user intervention in partition process
    - Unlike autoPart, partitionGroup.qel is not generated
  - Partition placement on rack resources done by largest partitions first
  - Only creates a partitionAssign file in tmp
    - Can be used in compile -etOnly to move existing partitions
  - The advancedPart option can only be respected in designs that require 32 or more domains
    - Smaller-than-32-domain designs will be scheduled into 1 partition

# How to Enable PPC: Enable Parallel Partition Compiler (Step 2)

- Step 2: enable parallel partition compiler
  - `compilerOption -add {distributedMultiCore ON | OFF | <path>}`
  - If this compiler option is turned ON, compilation runs in PPC mode, and each partition is compiled in its own directory.
  - The directories are created in PARTITIONS subdirectory within the design directory. You can also specify a <path> where you want to create the directories for the partitions.



# How to Enable PPC: Optimization (Step 3)

- Step 3 (optional): Determines what optimization are performed after partition databases are merged
  - `compilerOption -add {PostMergeOpt <level> }`
  - Can increase compile time
  - <value> can be:
    - 0 – no post merge optimizations
    - 1 – netlist optimization only
    - 2 – critical path optimization only
    - 3 – both netlist and critical path optimization
  - Netlist optimizations can reduce design size
  - Critical path optimization can help reduce step count by moving logic across domain boundaries

# How to Enable PPC: Compile (Step 4)

- Step 4: Compile as usual, using compile or compileFind command
  - Compile will spawn jobs for partitions automatically using LSF
  - When spawned jobs finish, perform final steps including scheduling
- LSF options can be with LSF\_options or with qel command
  - LSF Option in qel file will be used only for PPC jobs

```
set PPC(LSFOptionsForPartitions) {-q linux -R  
"rusage[mem=100000] select[mem>=200000 && (OSREL=EE60  
|| OSREL=EE70) ]" -M 200000}
```

# How to Enable PPC: Use Models

## Manual partitioning example

```
emulatorConfiguration -add .....  
partitionGroup -add {dut.a.b part1}  
partitionGroup -add {dut.a.b.r* part2}  
partitionGroup -add {dut.a.c.* part3}  
partitionGroup -add {* part4}  
partitionAssign -add {part1 0-1}  
partitionAssign -add {part2 2-3}  
partitionAssign -add {part3 4}  
partitionAssign -add {part4 5}  
compilerOption -add {PostMergeOpt 3} #optional  
compilerOption -add {distributedMulticore ON}  
compileFind .....
```

## autoPart example

```
emulatorConfiguration -add .....  
autoPart  
compilerOption -add {distributedMulticore ON}  
compileFind .....
```

## advancedPart example

```
emulatorConfiguration -add .....  
compilerOption -add {advancedPart ON}  
compilerOption -add {distributedMulticore ON}  
compileFind .....
```

# PPC: Successful Compile

- All stages of a compile must complete successfully
- Things to look for in xe.msg file
  - Did the compile successfully schedule and is there a final speed number reported?

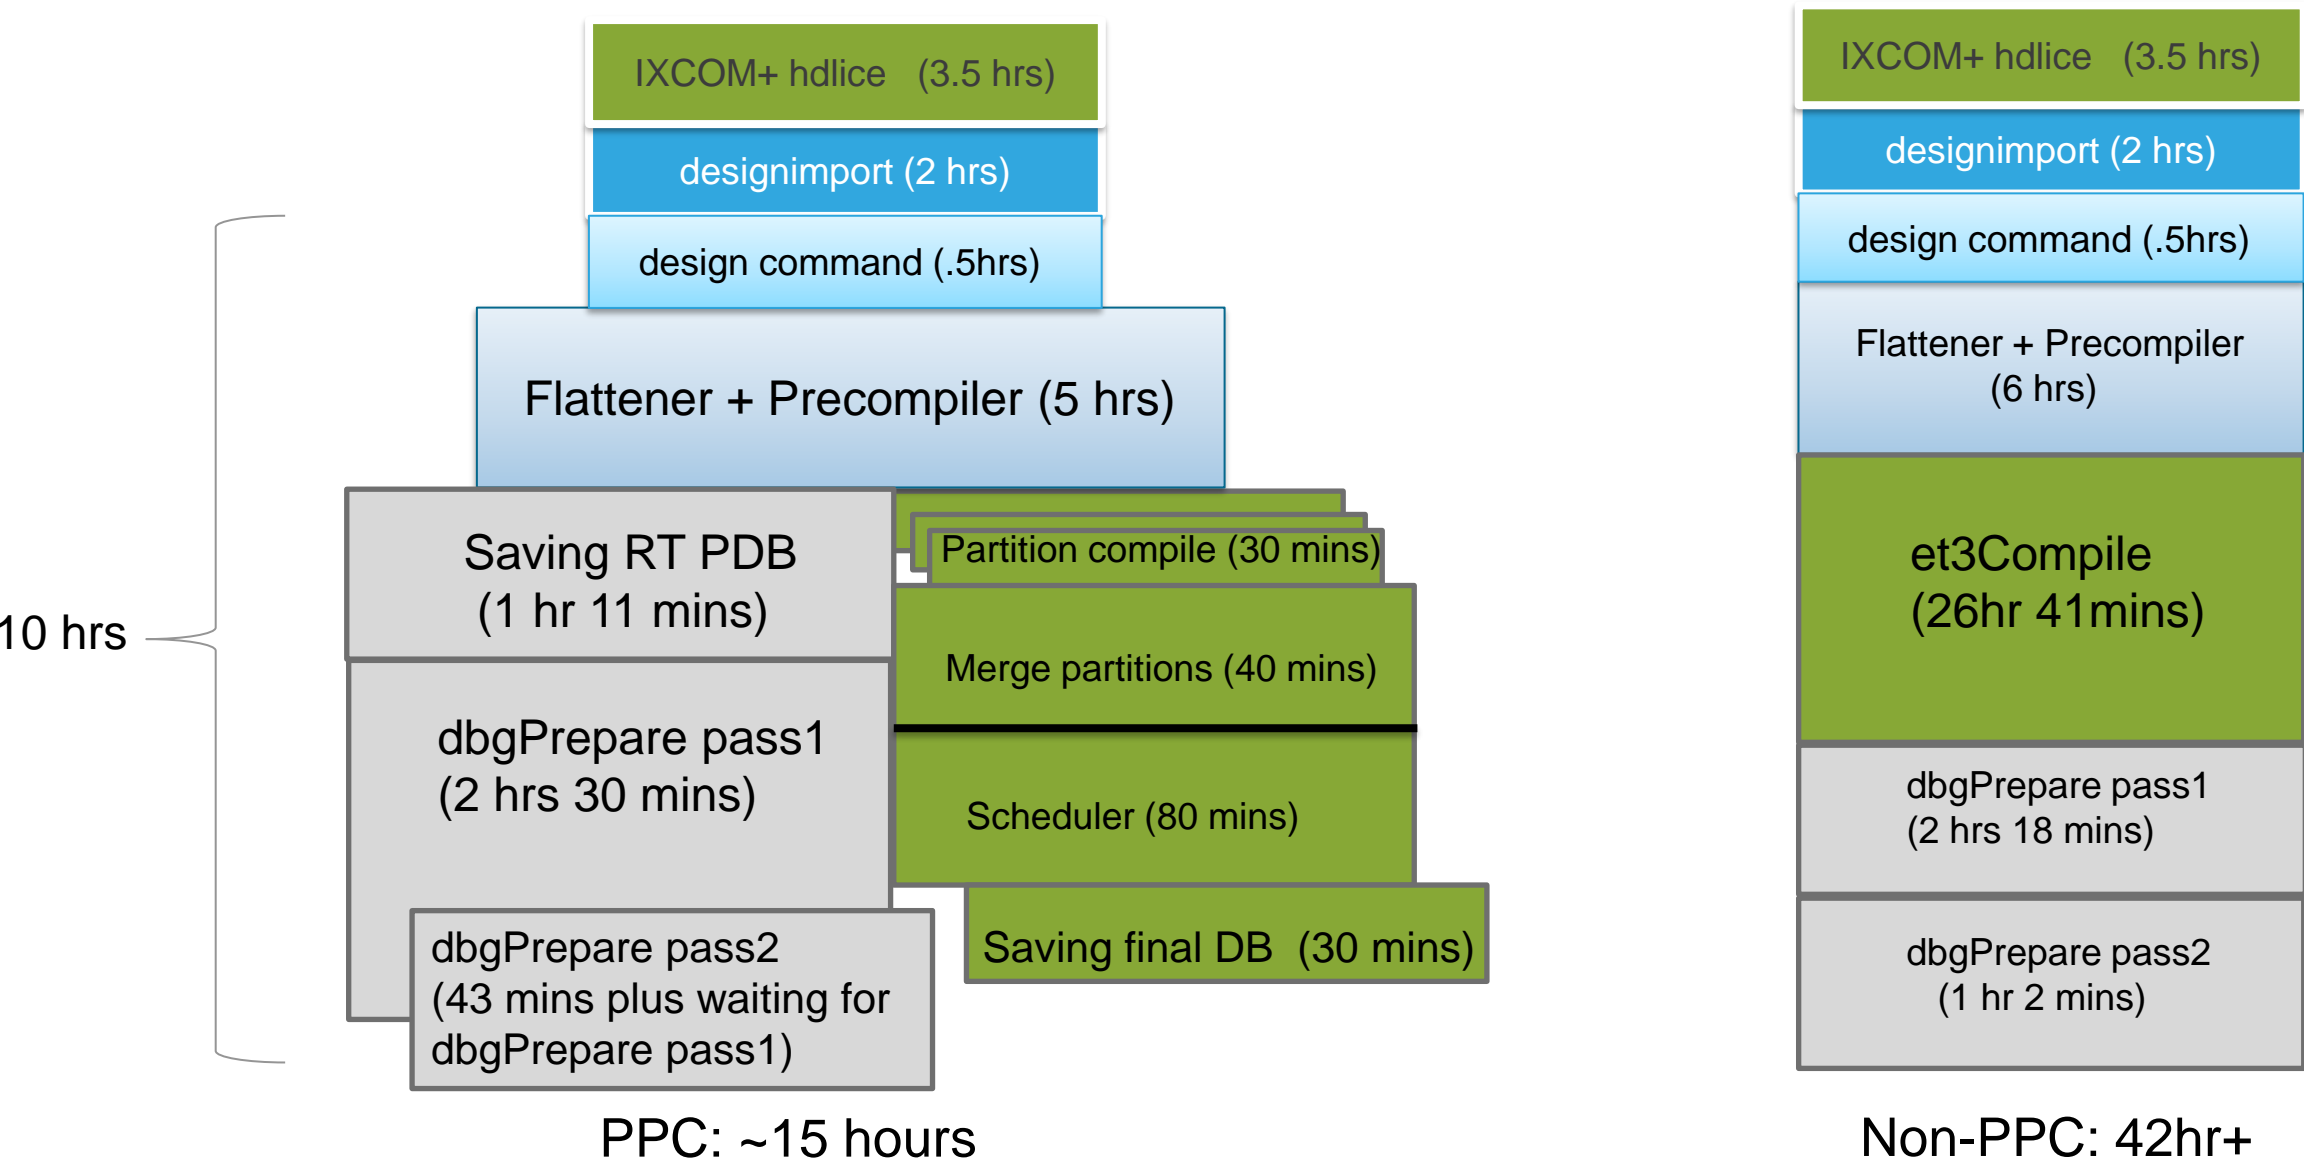
```
INFO (db2util-1067): This design is scheduled in 3728 steps.  
INFO (qt2dadb-1086): Maximum emulator operating speed is 138 kHz.
```

- Did dbgPrepare pass 1 and 2 complete? These 2 passes run in parallel but overlap manner, and both must complete

```
## dbgPrepare exits at 03/06/2019 03:19:51 Status: 0, PID: 87669  
## dbgPrepare Elapsed time: 7:38:25, U time: 33:46:37.56, Memory Usage: 4553.8G  
  
## dbgPrepare exits at 03/06/2019 03:19:57 Status: 0, PID: 132057  
## dbgPrepare Elapsed time: 2:11:52, U time: 29:43:54.54, Memory Usage: 833.4G
```

- Did xeCompile exit with no errors?

# Performance Comparison on a 3.78-billion-gate Design



# PPC: Considerations

- Designed for larger designs, in excess of 500M gates
- MUST have LSF or similar load-sharing facility available
  - Machines in LSF queues must have adequate resources - 1GB memory for 1 million gates
  - Limited error recovery for LSF compile is supported, will restart failed partitions
    - `compile -recoverPPC`
- If using manual partitioning
  - Recommended number of partitions is between 10 – 30.
  - Partition should be of equal size so they complete compile in equal time
- Refer to VXE User Guide on how to convert the process if choosing to use other load-sharing facility

# PPC: Considerations

- Optimal performance is achieved by NOT running precompile command
  - Specify all precompile options using `precompileOption –add` command in qel script
  - Then only use the compile command , this will automatically call precompile
- Optimizations are performed on partitions but they may not result in same optimization performed on the design as a whole
- Expectation is for 20% increase in emulator resources from sequential compiles
- Uneven distribution of memory in partitions can cause compile failure

# PPC: Compile Failure

- PPC compile failure will exit the compile
  - Other running partition compile jobs are terminated and xeCompile is aborted
- Logs for LSF / NC jobs for PPC are in tmp directory
  - Naming is reportFromPartition<partNumber>.log
- Partition compile error can be related to LSF issues. For example:
  - LSF queue / machine specification
    - wrong queue, machines resource mismatch, i.e. not enough memory
  - Machine's workload causes timeout
- Can restart compile and recover from PPC errors
  - **compile -recoverPPC**
    - Will check for PPC job error, recompile only those jobs and continue.



# Post PPC Issues

- If final scheduling does fail due to lack of resources

- Can retry with existing partition data
  - Recompile partitions, and re-do final scheduling
  - New seeds may alter how logic is aligned in partitions
- Use the xeCompile commands to retry

```
design <DesLib> <topCell>  
compilerOption -add {PostMergeOpt 2} # OPTIONAL  
compile -etOnly
```

- For IXCOM compile, can use -target xecompile:skipimport alternatively
  - Create compile.qel with the above commands
  - IXCOM starts compile at xeCompile stage, skips import and uses compile.qel

# Post PPC issues

- dbgPrepare can be a time consuming compile stage in PPC compile
  - dbgPrepare creates databases for debugging features at runtime
- dbgPrepare is done in two stages
  - pass1 takes more time and more memory
  - Pass2 will start concurrently but must wait for pass1 to complete
  - Memory footprint can be sum of both passes:

```
## dbgPrepare exits at 04/29/2019 02:51:34 Status: 0, PID: 168252
## dbgPrepare Elapsed time: 2:19:58, U time: 52:15:15.08, Memory Usage: 3028.4G
## dbgPrepare exits at 04/29/2019 03:38:43 Status: 0, PID: 168156
## dbgPrepare Elapsed time: 1:40:15, U time: 16:12:12.44, Memory Usage: 717.9G
```

- To compile single dbgPrepare pass
  - `dbgPrepare hlddb xcva_top -pass1 -mc`
  - `dbgPrepare hlddb xcva_top -pass2`
- To compile pass1 & pass2:
  - `dbgPrepare hlddb xcva_top`

# Outline

**Parallel Netlist (PNL) for IXCOM flow**

**Parallel Invocation of HDL-ICE Analyzer**


**Parallel Synthesis**

**Parallel Partition Compiler (PPC)**

**Modular Compilation Flow (MC)**

# Modular Compilation: Highlights

- Targeted to handle large designs without the need for powerful machines
- Better compilation turnaround time and lower memory footprints
- Distributes most of the compilation steps to compute farms like LSF or NC
  - A new command called diCompile controls distributed compilation.

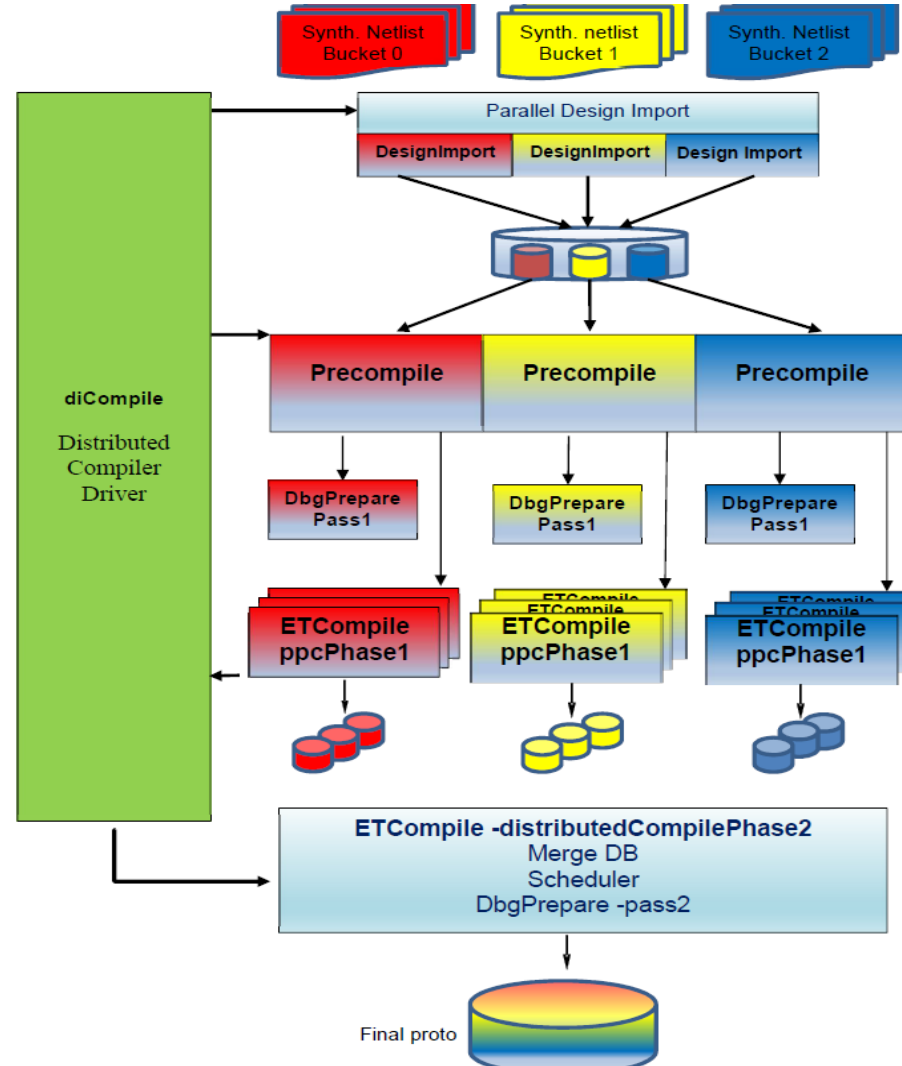


- For a 36-board design, MC is 2x faster than PPC
- For a 72-board design, MC is 2.7x faster than PPC
- For a 108-board design, MC is 4.5x faster than PPC

# Modular Compilation: Flow Chart

Synchronization points via diCompile

- At the end of parallel designImport
- During Precompile (through handshake)
- At the end of all ETCompile-ppcPhase1 jobs
- At the beginning of dbgPrepare-pass2 (wait for pass1 jobs to finish)



{advancedPart ON} is turned on by default and cannot be turned OFF.

{distributedMulticore ON} has no effect.

ETCompile -distributedCompile phase2 is launched on local host and multi threaded

# Modular Compilation: Preparation

- Design partitioning
  - User needs to partition the design into buckets
  - Each bucket should be big and logically distinct
  - In simple words, modules from one bucket should not instantiate those from another bucket
  - How big can each bucket be? What's the size of each bucket?
    - Sample DUT size: 10 Billion gates
    - Bucket size can be in the range of 500 Million gates to 5 Billion gates
    - Best MC performance is achieved by balanced distribution of modules into buckets
  - What if a user doesn't have any design knowledge? How to split the design?
    - Use gateCount utility and find the gate count info per module: `gateCount <libName> <cellName>`
- bucketModules.lst (file name is case sensitive)
  - Needed for providing user defined bucketing info
  - User creates this file with module names only. No instance names allowed
  - Buckets are internally numbered 1 ... N
  - Bucket 0 is implicit to contain design top and anything not included in user defined buckets
  - Bucket numbering is not based on the order of the listed modules in bucketModules.lst

# Modular Compilation: Puzzle#1

- Sample design top snippet

```
module modA;  
  ..  
  modG instG();  
  ..  
endmodule  
  
module modB  
  ..  
  modF instF();  
  ..  
endmodule  
  
module modC  
  ..  
  modD instD();  
  ..  
endmodule  
  
module modD  
  ..  
  modH instH();  
  ..  
endmodule
```

- Sample bucketModules.lst:  
modA  
modB  
modC  
modD
- Is the above bucketModules.lst appropriate?

# Modular Compilation: Answer to Puzzle#1

- Sample design top snippet

```
module modA;  
  ..  
  modG instG();  
  ..  
endmodule  
  
module modB  
  ..  
  modF instF();  
  ..  
endmodule  
  
module modC  
  ..  
  modD instD();  
  ..  
endmodule  
  
module modD  
  ..  
  modH instH();  
  ..  
endmodule
```

- Sample bucketModules.lst:  
modA  
modB  
modC  
modD
- Is the above bucketModules.lst appropriate?
  - No. modD is instantiated inside modC, so it should not be included in the list.



# Modular Compilation: Preparation (Continued)

- Ensure compute farm job management SW is installed (like LSF or NC)
- Provide required compute farm job options
  - Instead of using LSF\_options file as is done in non-modular compilation, modular compilation relies on a JSON configuration file named computeFarm.json that must be accessible in the design directory.
  - Unlike LSF\_options, computeFarm.json is mandatory for MC flow
  - Add “set MC(computeFarm) auto” in compile-time qel file
  - Precompile (aka bucket) jobs are launched with a single job ID
  - All precompile jobs are launched together when compute farm resources are available

# Modular Compilation: Preparation (Continued)

- Sample computeFarm.json for LSF

```
{
  "farmName" : "LSF",
  "farmOptions" : {
    "Partitions" : "-M 150000 -R \" select[OSNAME==Linux &&
(OSSTYPE==CL&&(OSMJR==6)) || (OSSTYPE==RL&&(OSMJR==5 || OSMJR==6 || OSMJR==7)) || (OSSTYPE==
SL&&(OSMJR==11))] rusage[mem=40000]\" -q secondary",
    "Buckets" : "-M 150000 -R \" select[OSNAME==Linux &&
(OSSTYPE==CL&&(OSMJR==6)) || (OSSTYPE==RL&&(OSMJR==5 || OSMJR==6 || OSMJR==7)) || (OSSTYPE==
SL&&(OSMJR==11))] rusage[mem=40000]\" -q secondary",
    "Import" : "-M 200000 -R \" select[OSNAME==Linux &&
(OSSTYPE==CL&&(OSMJR==6)) || (OSSTYPE==RL&&(OSMJR==5 || OSMJR==6 || OSMJR==7)) || (OSSTYPE==
SL&&(OSMJR==11))] rusage[mem=50000]\" -q secondary" ,
    "dbgPrepare" : "-M 150000 -R \" select[OSNAME==Linux &&
(OSSTYPE==CL&&(OSMJR==6)) || (OSSTYPE==RL&&(OSMJR==5 || OSMJR==6 || OSMJR==7)) || (OSSTYPE==
SL&&(OSMJR==11))] rusage[mem=40000]\" -q secondary"
  }
}
```

- Sample computeFarm.json for NC

```
{
  "farmName" : "NC",
  "farmOptions" : {
    "Partitions" : "-D",
    "Buckets" : "-D",
    "Import" : "-D" ,
    "dbgPrepare" : "-D"
  }
}
```

# Modular Compilation: Puzzle#2

- Where can I find the number of emulator domains used for each MC buckets?

# Modular Compilation: Answer to Puzzle#2

- Where can I find the number of emulator domains used for each MC buckets?
  - This information is in DISTRIBUTED/\*/dbFiles/<design>.et3config files:

```
+ CABLE_5 23 J23 30 J13
+ CABLE_5 23 J26 24 J36
* The configuration consists of 57 domains of 36 boards, 288 cables, 0 HDDC cable
DISTRIBUTED/0/dbFiles/xcva_top.et3config
+ CABLE_5 23 J23 30 J13
+ CABLE_5 23 J26 24 J36
* The configuration consists of 33 domains of 36 boards, 288 cables, 0 HDDC cable
DISTRIBUTED/1/dbFiles/xcva_top.et3config
+ CABLE_5 23 J23 30 J13
+ CABLE_5 23 J26 24 J36
* The configuration consists of 33 domains of 36 boards, 288 cables, 0 HDDC cable
DISTRIBUTED/2/dbFiles/xcva_top.et3config
+ CABLE_5 23 J23 30 J13
+ CABLE_5 23 J26 24 J36
* The configuration consists of 33 domains of 36 boards, 288 cables, 0 HDDC cable
DISTRIBUTED/3/dbFiles/xcva_top.et3config
```

# Modular Compilation: IXCOM Mode Mandatory Switches

- Enable parallel design import using IXCOM option -parallelNL
  - `ixcom -parallelNL <other required ixcom options>`
- Enable MC flow by invoking IXCOM using “-xeCompileArgs -mcBucket”
  - `ixcom -parallelNL -xeCompileArgs “-mcBucket”`
  - bucketModules.lst should be in the compile directory
  - bucketModules.lst cannot be left empty

# Modular Compilation: ICE Mode Commands

- MC flow in ICE mode starts from using generated QTDB libraries
  - User needs to generate the QTDB libraries
- User is expected to update compile.qel with below MC commands
  - *compilerOption -add {distributedLibraries ON}*
    - Enables Modular Compilation for ICE flow
  - *setModularCompile*
    - Checks for user errors before distributed compilation starts
  - *set MC(computeFarm) auto*
    - Required to use the json based compute farm framework
  - *diCompile*
    - Launches the distributed compilation

# Modular Compilation: Puzzle #3

- Is the number of PPC autopartitions under each MC bucket fixed?

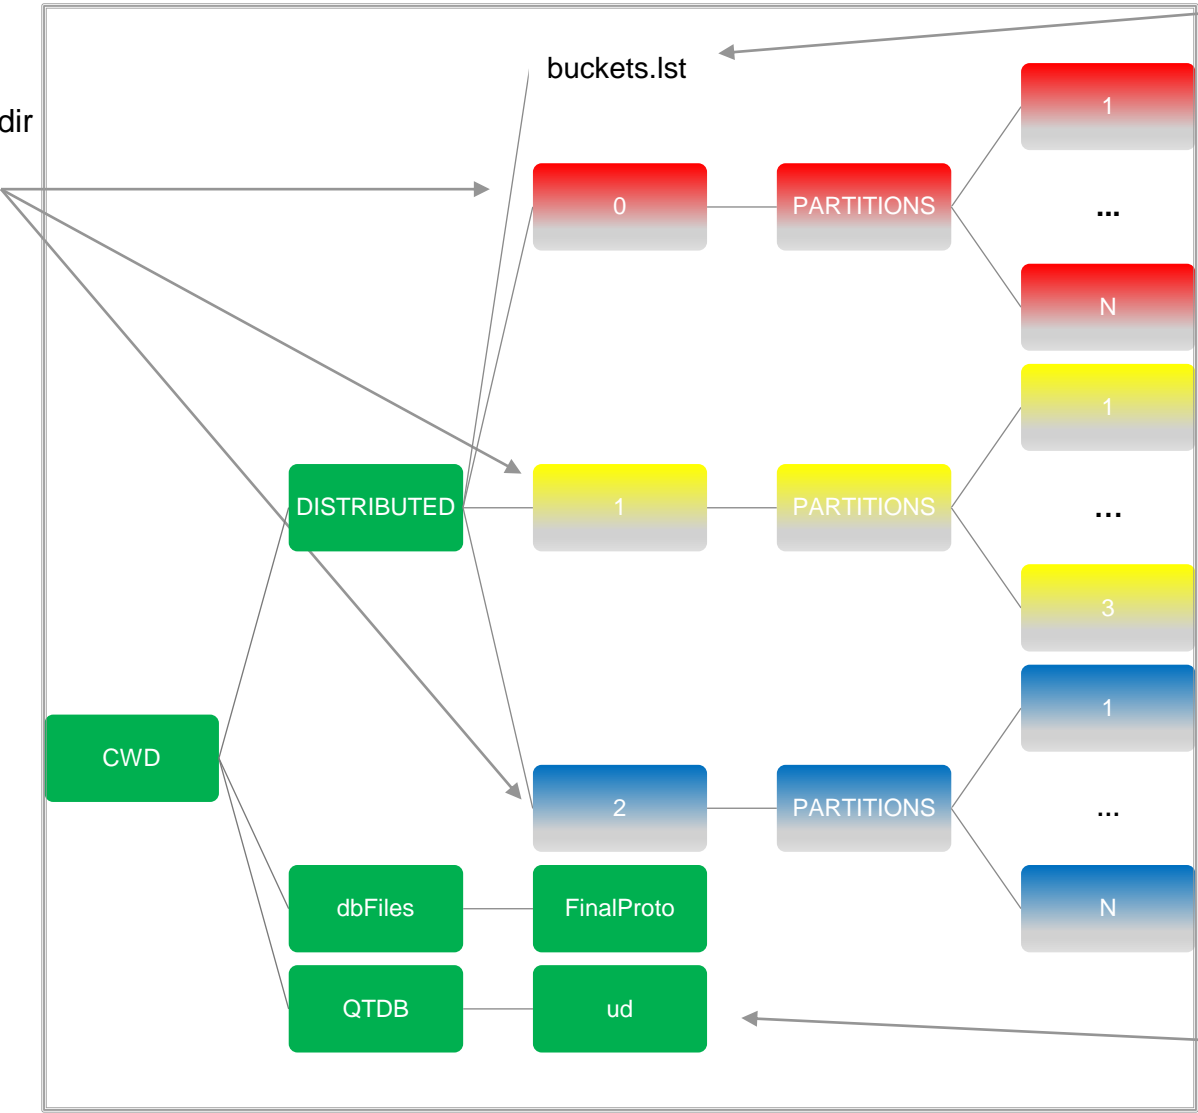
# Modular Compilation: Answer to Puzzle #3

- Is the number of PPC partitions under each MC bucket fixed?
  - No. They'll vary based on the size of the MC bucket if tool does the partitioning
    - PPC is turned on by default with {advancedPart ON}
  - Yes, if user creates the partitions manually by using below command
    - compilerOption -add {advancedPart N} ; N = 1,2...



# Modular Compilation: Generated Directories

In this scenario, user created 2 bkts + bkt0  
Sub dir for each bkt is under DISTRIBUTED dir



This file has hierarchical instance names for each bucket module listed in bucketModules.lst

Has user-defined precompile and compile information

# Modular Compilation: Puzzle #4

- How does a user know what is inside each bucket?

# Modular Compilation: Answer to Puzzle #4

- How does a user know what is inside each bucket?
  - From DISTRIBUTED/buckets.lst
  - Or, use the run-time debug command
    - `debug . -listbuckets`
    - Bkt0 is special and will be displayed as {}

# Modular Compilation: Using Compile Find Capability

- MC supports compile find capability provided by diCompileFind command
  - Multiple compile trials for larger designs to improve success rate
  - Faster compile time and reduced memory footprint compared to PPC
  - Use same compute resources as normal MC compile
- Main diCompileFind Options
  - -parallel: run all trials in parallel, and keep the solution with best step count
  - -serial: run all trials sequentially, and keep the solution with best step count
  - -serial\_best: run trials sequentially and aborting any trial which is not an improvement
  - -serial\_first: run all trials sequentially and stops at the first successful compilation
  - Options above are MUTUALLY EXCLUSIVE, in which the default is -serial\_best
  - Running all diCompileFind trials across compute farm can cause issues: disk usage for multiple trials can be large

# Modular Compilation: Full Debug Visibility

- The default option offers full debug visibility during a run-time session
  - Gives xeDebug access to the entire design PDB
  - Each bucket's PDB is opened by a PDB remote server
  - PDB server can run on local host, list of specified machines or across LSF farm
  - Complete design waveform can be viewed
- Use existing debug command
  - debug <path>

# Modular Compilation: Full Debug Visibility (Continued)

- Parameters or environment variable used to distribute across multiple machines
  - Default is to run all PDB server on local host
  - Bucket 0 always runs on the local host
- Two new parameters to control the launch of the PDB servers
  - `xeset pdb_server local | lsf | {<host1> <host2> ...}`
    - Specify the hardware to run remote servers
  - `xeset pdb_remote rsh| ssh`
    - Choose connection type to remote server hardware when using hosts
  - Must be done before the host command
- Values can also be set with environment variables
  - `setenv XE_PDB_SERVER local | lsf | "host1 host2"`
  - `setenv PDB_REMOTE ssh | rsh`

# Modular Compilation: Bucket-wise Visibility

- You also have the option to debug your design according to the defined buckets
- User needs to select which bucket to access at run time
  - By default, NO bucket is selected
  - Once selected the bucket cannot be changed in the same debug session
- Debug capabilities are limited to current bucket scope
  - KeepNet is required to access non-current bucket signals for commands like value, drtl, sdl, etc.
  - User can use memory commands on non-current bucket memory arrays without any keepNets

# Modular Compilation: Bucket-wise Visibility (Continued)

- Ways to choose bucket at run time
  - 1. Using debug command in run script
    - `debug . -bucket <leaf/bkt module inst name | bucket_number | none >`
    - Recommended to use leaf module instances rather than bucket number:
      - `debug . -bucket 2`
      - `debug . -bucket designTop.Inst2` *# Best approach*
      - `debug . -bucket none` *# Regression mode*
    - Bucket number may not remain the same when design or VXE build changes
  - 2. Using XE\_DEBUG\_ARGS env variable
    - best suitable for xrun mode
    - `setenv XE_DEBUG_ARGS ". -bucket 2"`
    - `env XE_DEBUG_ARGS=". -bucket 0"`
  - 3. Using xeDebug command
    - `xeDebug -input "@debug . -bucket ..." <other xeDebug options>`
- Regression mode
  - Ideal for quick validation with limited debug capabilities



# Modular Compilation: Bucket-wise Visibility (Continued)

- Online bucket mode
  - Only signals of current bucket can be seen in the waveform when generated online
  - Non-current bucket signals that are preserved by keepNets can be seen in waveform only when they are streaming probed
  - A single PHY directory is dumped
    - All the buckets 1 to N sub directories will be seen underneath PHY directory
- Regression mode
  - Regular probe command is not supported; thus, *shm/fsdb waveform cannot be generated during this online mode*
  - Dumping of the PHY file is the same as online bucket mode
- Offline mode
  - Whether PHY is generated using regression mode or bucket mode, the offline mode flow remains same
  - User can generate the waveform of any bucket (1 to N) by starting an offline debug session with `–bucket <N>`, add probes from this bucket, and then dump the waveform

# Modular Compilation: Restrictions

- Modular compilation does not support:
  - ATB constructs
  - UPF, CPF, and Hardware-Assisted Weighted Toggle Count (HW-WTC) commands.
    - Note that Modular Compilation supports Fast Toggle Count (FTC).
  - DPA and TCA operations
  - External netlists
  - RTL database
  - Executing precompile separately from compile
  - Following database commands:
    - db
    - moduleNet
    - getDatabaseInfo
  - partitionGroup and partitionAssign commands
  - Combinational loops involving more than one bucket, if detected before scheduling.
    - Such cases are reported in advance, so that the loop can be broken manually or the bucketing strategy modified as required.

# Modular Compilation: Restrictions (Continued)

- In addition to the above limitations, at times, when the simulator time is reset to 0, DCC waveform mismatch can be observed between InfiniTrace Prepare and Observe sessions.
  - This issue appears intermittently for modular-compiled designs. To workaround the issue, specify the time without the trigger keyword
    - Example: `infiniTrace -goto 1225`.

**cā dence®**

© 2020 Cadence Design Systems, Inc. All rights reserved worldwide. Cadence, the Cadence logo and the other Cadence marks found at [www.cadence.com/go/trademarks](http://www.cadence.com/go/trademarks) are trademarks or registered trademarks of Cadence Design Systems, Inc. All other trademarks are the property of their respective holders.